

Learning to drive based on multiple sensor cues in The Open Racing Car Simulator (TORCS)

Tim Bicker and Nimar Blume

Abstract—To implement an autonomous driver in The Open Racing Car Simulator (TORCS) we use a combination of a deep neural network (DNN) and a spiking neural network (SNN) based on multiple sensor cues. Specifically, the DNN predicts the current car displacement and angle relative to the road centre from a driver's view image. Based on the two values and an additional multi dimension range finder sensor a SNN generates driving commands for the car. Subsequently, the driving performance is evaluated on unseen tracks.

Index Terms—deep learning, TORCS, convolutional neural network, spiking neural network, autonomous driving

I. INTRODUCTION

AUTONOMOUS driving is currently a controversial subject, especially regarding feasibility and performance. To explore the details of autonomous driving, an abstract prototype is built focussing on the implementation of inferring location information from images and implementing a neural controller to drive a car. In this work we use The Open Racing Car Simulator (TORCS) to simulate a car on a race track and collect relevant sensor data. The goal is to use multiple sensor cues to drive the car safely around the defined test tracks. Further, we use the Robot Operating System (ROS) TODO: ADD CITATION as an interface between TORCS and the controller TODO: ADD ROS TORCS CITATION. The following signals are used:

- 1) range finder
- 2) driver point of view (POV) image
- 3) displacement (distance from road centre to car centre)
- 4) angle (rotation angle between car and road)

Then, a deep neural network will be trained on the ground truth data to infer the angle and displacement from the provided POV image. Subsequently, a controller based on spiking neural networks (SNN) is trained to generate driving commands based on range finder, displacement and angle sensor inputs. Finally, the controller and the DNN are connected. So the final model works only based on the POV image and the range finder data. An overview of the overall structure can be seen in figure XX.

TODO: model structure overview picture as in presentation

II. A DEEP NEURAL NETWORK FOR REGRESSION

The implementation of the deep neural network is done in python version 3.6.3. To facilitate development, the library

Authors: Tim Bicker (12345678, tim.bicker@tum.de) and Nimar Blume (03638934, nimar.blume@tum.de) **Course:** Practical course: Computational Neuro Engineering Winter Semester 2017/2018 **Submitted:** January 9, 2018 **Supervisor:** Florian Mirus, Neuroscientific System Theory (Prof. Dr. Jörg Conradt), Technische Universität München, Arcisstraße 21, 80333 München, Germany.

keras [5] is used with the backend tensorflow [3]. To manipulate image data openCV3 [4] is used and finally to load and manipulate data the library numpy [12] is utilised.

A. The testing deep neural network

To be able to chose hyperparametres for the convolutional neural network (CNN) at an early point, a basic CNN network architecture was chosen which has been proven before. The criteria for choosing the network architecture was, that it has to provide reasonable performance while being quick to train. The focus was rather on fast training performance, as the available resources are limited to us. Thus, after studying the architectures' training performances as seen here [6], AlexNet [8] was chosen as testing DNN.

AlexNet is a network designed for image classification tasks, such as the ImageNet challenge. Therefore, the last layer of AlexNet was altered to use it for multidimensional regression problems such as inferring the angle and displacement from an image. To achieve that, the number of output neurons of the last fully connected layer (FCL) was reduced from 1000 to 2, as there are two numbers to determine. Furthermore, the last FCL used a rectified linear unit (ReLU) as activation function.

$$f(x) = \max(x, 0) \quad (1)$$

To prevent the activation function from cropping the output to values greater than zero, a linear activation function is used instead: $f(x) = x$.

B. Data set splitting

TORCS provides 19 tracks of which all are relevant. Therefore, the recorded data set is split into three parts:

- 1) Training set
- 2) Validation set
- 3) Test set

First, the two tracks TODO:FIX TEST TRACKS were determined to be used for testing later on. As the goal is to train a general model, a prerequisite is that data recorded on either of the two test tracks is not used during training. In total XXX data points were recorded on the two test tracks.

The remaining tracks are used to train the DNN and to determine its hyperparametres. Therefore, the data set is randomly split into train and val (validation) at a ratio of 90% to 10%. The validation set data itself was thus not used for training, but it is recorded on tracks which are included in the training tracks.

C. Input images

The image data is acquired from TORCS [13] using ROS via the TORCS-ROS node [9].

1) *Choosing a camera angle:* TORCS provides several camera perspectives:

- 1) Driver's view with hood
- 2) Driver's view without hood
- 3) Third person perspective: far
- 4) Third person perspective: close

All perspectives are evaluated based on the DNN described in subsection II-A and as final perspective the driver's view without hood is chosen. The basis for that decision can be seen in Figure 1, which shows the mean absolute error for each view, with the driver's view without hood being the lowest. TODO: INCLUDE EXAMPLE IMAGES



Figure 1. Mean absolute error of the same DNN trained with multiple camera angles

2) *Choosing the image size:* The images are provided by TORCS ROS [9] at a rate of 10 frames per second (fps) at a resolution of 640 px \times 480 px. Because that image size is too large to train a reasonably deep network in a reasonable time, the images are down-scaled prior to training as well as in the final application. To determine the image providing the best result, four different sizes were evaluated with the DNN described in subsection II-A:

- 1) 320 px \times 240 px
- 2) 160 px \times 120 px
- 3) 80 px \times 60 px

First, the training images were collected from TORCS ROS at a resolution of 640 px \times 480 px. Upon loading the images, openCV based downscaling was applied using the bilinear interpolation algorithm. Second, the testing DNN was trained with the images at the mentioned resolutions and the mean absolute error for the validation set was recorded, which is visualised in Figure 2.

Therefore, the image size 80 px \times 60 px is used for the DNN as it provides the smallest mean absolute error and additionally reduces the training time due to the small image size.

D. Choosing the optimiser

A plethora of different optimisers are available for use in the backpropagation step when training DNNs. The largest differences are in their abilities to exit large local maxima, training performance and accuracy. As the currently used DNN is manageable in size, the training performance is not yet a critical property. Therefore, the optimiser was selected based on its ability to minimise the mean absolute error of the validation set. As illustrated in Figure 3, the adamax optimiser provides the best result and is thus chosen for future




Figure 2. Mean absolute error of the same DNN trained with multiple image resolutions

application. adamax is similar to the adam optimiser while being more stable in certain conditions, see [7].




Figure 3. Optimiser comparison in the test DNN using with an image size of 80 px \times 60 px

E. Building a deep neural network with keras

Keras is a python library hello keras [11]

III. NENGO CONTROLLER

The controller of the car is designed based on Spiking Neural Networks (SNN) which are implemented in Nengo [2], a python framework for building large-scale neural systems that is based on the Neural Engineering Framework (NEF) [10].

The controller is divided in several modular parts: acceleration, braking, steering, gear changing and clutching. We choose the modular approach, because it allows to test different modules independently from each other and also mix hard coded solutions with learned ones.

A. Nengo controller design

In the default driver's source code, the driving modules are based on different sensor inputs. The steering module is a function based on the car speed v_x , the lateral displacement

d_{lat} and the rotation angle α . Both, the acceleration and the braking module are based on three range sensor signals -5 , 0 and 5 degrees as well as v_x . We chose to implement the gear changing as a hard coded solution without any learning involved, since it is a simple if-else condition. It is further possible to neglect the clutching module, because it did not show to make any difference for driver performance even though it is implemented in the default driver. The final model architecture can be seen in Figure 4.




Figure 4. Nengo architecture with different modules. The input network contains an input node that has subscribed to the ROS topics and contains the CNN for inference of rotation angle and lateral displacement. The steer, accelerate and brake module contain one ensemble each. The output network contains an output node that publishes to the ROS topics. The gear module is not shown here.

B. Controller Training

In this work we chose a supervised learning approach for training the controller. Nengo offers two types of learning strategies: offline and online learning. With Nengo deep learning [1], an extension to implement classic deep learning methods exists as well.

a) *offline learning*: can be used for the classical supervised learning approach. Feature data and corresponding labels are given and Nengo uses a least squares approximation to solve for the neuron weights [10].

b) *online learning*: iteratively improves the initial neuron configuration in a supervised manner during runtime. However, this requires that for every timestep during runtime a goal value has to be known for Nengo to solve for it. In our problem this is not the case. Therefore online learning is not further considered in this work.

c) *Nengo deep learning*: allows to train classical deep learning methods e.g. feedforward neural networks with back-propagation and then implement the learned nets into Nengo. This work focuses on a mixture of hard coded functionality and offline learning.

C. Evaluation metric

To evaluate the controller's performance, we used the following two metrics: Firstly, we investigate if the controller is able to finish one lap in a reasonable time and how fast he is compared to the default driver. Secondly, we judge the overall

robustness of the controller. For this we observe the actual controller during driving and evaluated how stable he drives, e.g. how close he actually drives to the edge of the street.

D. Data Sources

a) *TORCS-ROS default driver*: The first and easiest to acquire is the default TORCS-ROS driver [9]. The advantage of this driver is that he drives a maximum speed of $149 \text{ km} \cdot \text{h}^{-1}$ and drives very carefully. Which means, that he always tries to stick to the middle lane of the road and brakes heavily when he comes close to a turn. We believe that this is an easy to learn and robust driving style. However, this driver is very slow, as can be seen in Table I.

b) *TORCS drivers*: Furthermore, there are several drivers that are implemented in TORCS. Those drivers have no speed limit and their driving style approximates that of a real race. The cars drive very close to the edge of the road and drive very fast through turns. We think that this is a hard to learn and error-prone driving style, because small errors lead to disastrous effects.

c) *driving manually*: Finally, there is the option to drive manually. As it turns out this is really difficult with a keyboard. Because fractions of seconds on the left and right arrows lead to very strong steering behavior and almost every time lead to a full turn-around or a far deviation of the track. This leads to a manual driving style that comes close to the default driver: very slow through turns and sticking close to the road center. However, because it is possible to drive as fast as one wants and in general one achieves a smoother driving of the car, we achieve a faster time than the default driver.

Because of the above mentioned characteristics of the differ-

TORCS-ROS default	85s
TORCS	43s
manual	69s

Table I

THIS TABLE COMPARES THE LAP TIMES OF DIFFERENT DRIVER TYPES.

ent driver styles, it does make sense to first learn the TORCS-ROS default driver, because he seems to be the most robust one and data generation is simple.

IV. EXPERIMENTS AND EVALUATION

A. Performance comparison between inferring angles vs. inferring displacement

For the final application, the same model is used to infer the car's displacement and its angle. During training, the loss function (MSE) and the metrics (MAE) are continually evaluating the consolidated training performance. However, to determine the performance of predicting the individual values, two models were trained on each either predicting the angle or predicting the displacement of the car. Figure 5 shows the difference between angle and displacement prediction is shown. In training and validation set, the value of the displacement is in the range of -0.9 to 1.1 and the angle is in the range of -0.9 to 1.1 . Thus, it can be concluded that the displacement prediction performance especially still needs to be worked on.




Figure 5. Mean absolute error of two test networks trained only to predict either the displacement or angle of the car

B. Controller performance

Controller performance is influenced by different factors. Data composition, model capacity and forward propagation time are discussed in the following.

In order to define the best model capacity, it is important to generalize well. The final model should not overfit and neither underfit. Besides the neuron sizes for the different modules and signal encodings, also neuron composition is of interest. It is possible that one module has a much more complex function and therefore needs a higher capacity to learn its corresponding function. During testing, however, we didn't see different module capacities to have a strong influence on the model. We increased the neuron sizes from a bad performing configuration until we hardly discovered any improvements on the generalization task. We could not reach overfitting with the dataset we used, because our hardware would eventually run out of physical memory.

When we collected track data for training the controller, we chose tracks with somewhat different shapes and turns in order to cover a wide variety of situations. We observed that small training sets would lead to overfitting of the model and for small model capacities to instable performance. Once we had enough data, the track composition did not seem to influence controller performance.

During the testphase we observed a large variance in the data. Even when we ran tracks twice with the same configurations we obtained different track times with up to 2s difference. Further the same models showed different behaviour in critical turns. We have not yet found out the reason for this, but we suppose it has something to do with the controller frequency, which is influenced by the kernel management and other processes running on the system. Because our setup was Ubuntu 16.04 on a virtual machine, it could be possible that a native host and therefore a more powerful system would show less variance.

As already mentioned it is important to keep the control frequency high. This is necessary because with increasing model capacity, the calculation time of the controller becomes too large which lead to an instable controller. Low frequencies lead to the controller signals being already outdated when they

arrive at the car and therefore controlling becomes instable. This is especially important on computers with older hardware, which, even with native operating systems, have problems to simulate TORCS, ROS and Nengo at the same time. We observed that for average controller iteration times larger than 12 ms the controller already becomes instable. One forward propagation of the Nengo controller without a DNN takes approximately between 1.5 and 3 ms. This requires the DNN to have a maximum forward propagation of 9 ms.

Another influence on controller performance is the way we acquire the signal. The forward propagation of the DNN in the input node (see Figure 4) leads to a delay until the calculated signal enters the SNN. As stated above, only the steering module is based on the DNN signals. In order to gain the best controller performance, we updated the range finder signals and the speed signals asynchronously. So even when the DNN was calculating the angle and displacement values, the Nengo input node still received updates from the ROS topics. Then, after the DNN had calculated the signals, we fed the now outdated angle and displacement values together with the most recent speed and range finder signals into the SNN.

As the final neuron sizes for the different modules we chose 2,000 neurons encoding each input signal. The driving modules consist of 2,400 neurons each. With smaller model capacities we obtained a less stable driver, while a larger model capacity did not show any improvements. Table II shows the results of the controller without the DNN, with and without 9 ms artificial interruption with asynchronous signal updates based on ground truth data input signals.

track	default driver	no interruption	interruption
cg track 2	1:58:15	1:59:76	2:05:84
wheel 2	4:21:45	4:26:80	4:46:51

Table II

THIS TABLE COMPARES THE LAP TIMES OF THE CONTROLLER WITH AND WITHOUT 9 MS ARTIFICIAL INTERRUPTION AGAINST THE DEFAULT DRIVER ON THE GENERALIZATION TASK WITH GROUND TRUTH DATA AS INPUT SIGNALS.

C. Controller and DNN

V. CONCLUSION

A. Outlook

In order to increase the controller performance it is possible to merge the accelerating and braking module into one module, so the controller can either accelerate or brake. Nengo Deep Learning is an interesting approach to see if controller performance could be improved. Further, a quantitative measurement of controller stability could help in finding the optimal controller capacity. The training data can be changed so that the driver is based on more range finder signals in order to steer earlier when approaching a turn. Additionally a lower speed limit seems promising.

REFERENCES

- [1] Deep learning integratino for nengo. Accessed: 2017-12-13. 3
- [2] The nengo neural simulator. Accessed: 2017-12-10. 2

- [3] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wat-tenberg, M. Wicke, Y. Yu, and X. Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. Software available from tensorflow.org. 1
- [4] G. Bradski. The OpenCV Library. *Dr. Dobb's Journal of Software Tools*, 2000. 1
- [5] F. Chollet et al. Keras. <https://github.com/fchollet/keras>, 2015. 1
- [6] J. Johnson. Cnn architecture benchmarks. <https://github.com/jcjohnson/cnn-benchmarks>, 2017. 1
- [7] D. P. Kingma and J. Ba. Adam: A method for stochastic optimization. *CoRR*, abs/1412.6980, 2014. 2
- [8] A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in neural information processing systems*, pages 1097–1105, 2012. 1
- [9] F. Mirus. torcs ros. https://github.com/fmirus/torcs_ros, 2017. 2, 3
- [10] T. C. Stewart. A technical overview of the neural engineering framework. Technical report, Centre for Theoretical Neuroscience, 2012. 2, 3
- [11] A. Toshev and C. Szegedy. Deeppose: Human pose estimation via deep neural networks. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2014. 2
- [12] S. van der Walt, S. C. Colbert, and G. Varoquaux. The numpy array: A structure for efficient numerical computation. *Computing in Science Engineering*, 13(2):22–30, March 2011. 1
- [13] B. Wymann, E. Espié, C. Guionneau, C. Dimitrakakis, R. Coulom, and A. Sumner. TORCS, The Open Racing Car Simulator. <http://www.torcs.org>, 2016. 2